

SHOULD A COMPANY BE ALLOWED TO SAY NO?

AI, the Military, and the Weekend Everything Changed

What if the government told your company to do something wrong?

What if saying no cost you everything?

What if millions of people noticed — and switched sides?

VOCABULARY: WHAT IS A CONSCIENCE?

conscience (n.):

Your inner sense of what is right and wrong — even when no one is watching and no rule forces you to act.



It is internal — not a law, not a boss, not a punishment.



It can tell you something is wrong even when it is technically legal.



Acting on it can cost you something — money, power, or safety.

Coming up: Anthropic CEO Dario Amodei used this exact word when refusing the Pentagon.

FOLLOWING RULES vs. ACTING ON CONSCIENCE

Following the Rules

- ▶ You do what the law says — nothing more
- ▶ You stop only when a rule forces you to
- ▶ Someone else decides what is right
- ▶ Example: A company stops polluting only after it gets fined

Acting on Conscience

- ▶ You do what you believe is right
- ▶ You stop even when no rule requires it
- ▶ You decide what is right — and accept the consequences
- ▶ Example: A company refuses a contract it finds morally wrong

When Anthropic said no to the Pentagon — which column did they fall into?

BEFORE WE GET TO THE STORY — THIS HAPPENED FIRST

Last weekend, something went viral. Here is what happened.

#1

Claude hit the #1 spot on the Apple App Store for the very first time — Sunday, March 1

Millions of people switched away from ChatGPT in one weekend.

But why?



Katy Perry Posted This



She posted a screenshot of the Claude Pro subscription page on X — with a red heart drawn around it. Millions of views. She was publicly siding with the company that stood up to the government.

#QuitGPT — people publicly canceled ChatGPT and switched to Claude to protest OpenAI's deal with the military.

#QUITGPT — WHEN POLICY BECOMES CULTURE

A technical policy fight became a mainstream moment. Here is why it spread.



Katy Perry's Post

A pop star with millions of followers publicly endorsing an AI app — because of ethics, not features. Her post reached millions of people who had never heard of Anthropic.

Specific caption details are reported — we present confirmed facts only.



Claude Hits #1

For the first time ever, Claude reached #1 on the Apple App Store on Sunday March 1. A direct result of the wave of users switching from ChatGPT after OpenAI's Pentagon deal went public.

Confirmed — podcast hosts mentioned this directly.



The Movement

#QuitGPT spread across social media. People used their subscription dollars as a protest vote. The message: if you make tools for military surveillance, we take our money elsewhere.

Consumer activism — app store rankings as protest tool.

Discussion: Is switching apps a meaningful form of protest? Can consumer choices change how companies behave?

MEET THE PLAYERS

Four groups. One big fight. Here is who you need to know.



Anthropic

The company that makes Claude. Founded to build AI safely. CEO: Dario Amodei.

→ **Said NO to the military's demands. Got punished.**



OpenAI

The company that makes ChatGPT. The biggest AI company in the world.

→ **Said YES to a deal with the military. Got the contract.**



The Pentagon

The US Department of Defense — America's military headquarters.

→ **Wanted to use AI tools for military operations — on their own terms.**



Sam Altman

CEO of OpenAI. On Thursday he said he agreed with Anthropic. By Friday he had signed a deal with the Pentagon.

→ **Reversed course in less than 24 hours.**

WHAT DOES THE MILITARY WANT AI FOR?

AI is not just chatbots. Here is why it matters to the military.



Analyzing huge amounts of data instantly

AI can scan millions of social media posts, text messages, and location records in seconds — something that would take humans years.



Making decisions in combat

Some military leaders want AI to help decide when and where to use weapons — without waiting for a human to give the final order.



Watching people

AI can track who people talk to, where they go, and what they post online. This can identify threats — or be used to spy on people who have done nothing wrong. It could be deployed domestically.

THE TWO THINGS ANTHROPIC REFUSED TO DO

These are called their "red lines" — limits they said they would NEVER cross, no matter what.



RED LINE 1: Watching Americans Without Permission

Imagine if the government used AI to read every social media post in the country — looking for people it doesn't like. Anthropic said: we won't build that.

REAL EXAMPLE:

Real example: The government is already scanning social media posts of people trying to move to the US — looking for posts that criticize the president.



RED LINE 2: Weapons That Kill Without a Human Decision

Imagine a drone that decides on its own to fire a missile — no person pressing a button. Anthropic said: we won't build that either.

REAL EXAMPLE:

These are sometimes called 'killer robots.' Most military experts agree they are dangerous — but there is no law yet that bans them.

BUT ISN'T THAT ALREADY ILLEGAL?

Great question. Here is the problem.

The Pentagon said: 'Don't worry — we'll only use AI for things that are legal.'

So what is the problem with that?



The US has no law about AI surveillance

There is no rule saying the government can't use AI to watch millions of Americans. 'Only legal things' is a very big loophole.



The government can legally buy your data

Companies sell data about where you go, what you buy, and what you post. The government can buy all of it legally — it is not technically called surveillance, but it works the same way.



No one can check what happens inside

Once Claude is inside the military's classified computers, Anthropic cannot see how it is being used. There is no referee.

WHAT DARIO AMODEI ACTUALLY SAID

Dario Amodei is the CEO of Anthropic. This is what he wrote publicly.



We cannot in good conscience accede to their request.

— Dario Amodei, CEO of Anthropic, February 26 2026

"good conscience"

He is saying this is about ethics — not just business. Conscience means your own inner sense of right and wrong.

"accede to"

This just means 'agree to' or 'give in to.' He is saying: we are NOT giving in, no matter what it costs us.

The price they paid:

The US government ordered every federal agency to STOP using Claude. Anthropic lost millions of dollars in government contracts.

48 HOURS THAT CHANGED EVERYTHING

Here is what happened in order — from Thursday morning to Saturday night.



The strange part: Sam Altman said Thursday morning he agreed with Anthropic — then 24 hours later signed the exact deal Anthropic had just refused.

SAM ALTMAN — A CLOSER LOOK

Here are the facts. You decide what to make of them.



Sam Altman
CEO, OpenAI



Fired by his own board (2023)

OpenAI's board fired him, saying he was 'not consistently candid.' Employees demanded he return. He was back within days — but his own colleagues said he misled them.



No one can read the contract

The deal OpenAI signed with the Pentagon is classified. Anthropic says they were offered similar language but it was full of loopholes. We cannot verify Sam's promises.



Sora and copyright

OpenAI trained its video AI Sora on copyrighted content — and made artists opt OUT rather than opt IN. Many creators only found out after the fact.



The AI bubble

Sam has raised hundreds of billions of dollars betting AI will change everything. Many economists say current AI revenue does not justify those valuations. He is the face of a very expensive bet.

THE 24-HOUR FLIP

Sam Altman said two very different things in less than 24 hours.



Thursday Morning



"We stand with Anthropic. OpenAI will NOT allow our models to be used for mass surveillance or autonomous weapons."

Sam Altman sends this to OpenAI employees privately.



Friday Night



"OpenAI has reached an agreement with the Pentagon to deploy our models in their classified network."

Sam Altman posts this publicly on X — with the same two limits Dario was just punished for defending.

THIS IS NOT THE FIRST TIME



BACKGROUND:

In the 1940s, US scientists built the world's first nuclear bomb (the Manhattan Project). The lead scientist — Robert Oppenheimer — later tried to limit how it was used. The government destroyed his career.



1940s — The Atom Bomb

What they built:

The world's most destructive weapon — for the US military during World War II.

What they said:

After the war, Oppenheimer said it was being used too recklessly and pushed for limits.

What happened:

The government revoked his security clearance and ended his career. He was labeled disloyal.



2026 — AI Models

What they are building:

The world's most powerful AI tools — now being sought by the military.

What they said:

Anthropic's CEO said the government's terms were unsafe and refused to sign.

What happened:

The government cut all contracts and called Anthropic a national security threat.

WHY ARE WE MAKING THIS COMPARISON?

**This is not just a metaphor.
Dario made this connection himself.**



The Book

Dario Amodei bought every single Anthropic employee a copy of a book about how the atom bomb was built in the 1940s. His message to his team: study this — because what we are building is just as powerful, and one day the government will come for it.



He believed this day was coming

Dario is not alone — many of the world's leading AI researchers believe AI will change the world as much as nuclear weapons did, maybe more. This is not hype. These are the people who actually build it, and they are genuinely worried.



This week, his prediction came true

The government came for Anthropic's technology — and demanded to use it on their own terms. Exactly what the book warned about.

THE BIG QUESTION

Should a company be allowed to say NO to the government?



The Government's Argument

In a democracy, elected leaders — not private companies — make decisions about national security. If a company doesn't like a policy, change it through voting — not by refusing.



Anthropic's Argument

There is no AI surveillance law yet. Someone has to draw a line. The people who build the technology understand better than anyone what it can do wrong.



The Harder Question

If OpenAI will build it anyway, did Anthropic's refusal actually protect anyone? Or did it just hand the contract to a company willing to say yes?

LET'S TALK — THEN WRITE

Class discussion — then pick one exit ticket question to answer in writing before you leave.

DISCUSSION

Was Anthropic right to say no, even if it cost them millions?

Does Sam Altman's 24-hour flip make him a traitor, a smart businessman, or something else?

Should the government be allowed to use AI to read everyone's social media — even to find criminals?

Millions switched apps as a protest. Does switching apps actually change anything?

EXIT TICKET

Pick ONE — write before you leave:

A Was Anthropic right? Give one reason.

B Should companies refuse government contracts? Yes or no — why?

C What surprised you most about this story?

D What question do you still have?

(Specific caption details about Katy Perry's post are reported — we presented confirmed facts only.)